


RQ2 cost: score and retrieve stage per system

Log axis: fastest to slowest spans about 59x, so a linear axis would collapse three of these rows onto one another.



Latency times the evaluated configuration: the full Dawn Whitmore store with the shared deterministic embedder on both the write and query paths. Nearest rank percentiles over 100 timed retrievals per system, wall clock on one machine. This is the one measurement in the run that is not deterministic, and the store holds a single scenario's memories, so read the ratio between systems rather than the absolute milliseconds.

run 20260817-160950 | commit 55e6533452c0 (dirty tree) | model stub | labels dawn-whitmore v1 | built by assets/build_figures.py

PRELIMINARY: stub model (echoes the player's line), deterministic lexical embedder, v1 single author labels, 10 queries. Every interval spans zero and no Holm corrected comparison is significant: read direction, not ranking.