


RQ2 memory poisoning: injections that reached retrieval

10 injection attacks per system (2 categories of 5); a bar counts the attacks whose planted memory entered the probe's top 5.



The other 10 attacks (role override, persona dissolution) are absent here by construction, not zero by measurement: they write nothing to the store, so no poison exists to retrieve. Drawing them as zero bars would claim a defended result where the architecture has nothing to defend. The measurement that establishes it is `probe_prompt_identical`, true for all 10 of them in every system and false for every injection.

run 20260817-160950 | commit 55e6533452c0 (dirty tree) | model stub | labels dawn-whitmore v1 | built by assets/build_figures.py

PRELIMINARY: stub model (echoes the player's line), deterministic lexical embedder, v1 single author labels, 10 queries. Every interval spans zero and no Holm corrected comparison is significant: read direction, not ranking.