


Bake-off: per-turn generation latency by model



lower is better; cloud arms include network time